

Gaussian

July 26, 2024

```
[5]: %pip install fredapi
```

```
Collecting fredapi
  Downloading fredapi-0.5.2-py3-none-any.whl.metadata (5.0 kB)
Collecting pandas (from fredapi)
  Downloading pandas-2.2.2-cp311-cp311-macosx_11_0_arm64.whl.metadata (19 kB)
Collecting numpy>=1.23.2 (from pandas->fredapi)
  Downloading numpy-2.0.1-cp311-cp311-macosx_14_0_arm64.whl.metadata (60 kB)
    60.9/60.9 kB
4.5 MB/s eta 0:00:00
Requirement already satisfied: python-dateutil>=2.8.2 in
/Users/neerajnamani/miniconda3/envs/py311/lib/python3.11/site-packages (from
pandas->fredapi) (2.8.2)
Collecting pytz>=2020.1 (from pandas->fredapi)
  Using cached pytz-2024.1-py2.py3-none-any.whl.metadata (22 kB)
Collecting tzdata>=2022.7 (from pandas->fredapi)
  Using cached tzdata-2024.1-py2.py3-none-any.whl.metadata (1.4 kB)
Requirement already satisfied: six>=1.5 in
/Users/neerajnamani/miniconda3/envs/py311/lib/python3.11/site-packages (from
python-dateutil>=2.8.2->pandas->fredapi) (1.16.0)
Downloading fredapi-0.5.2-py3-none-any.whl (11 kB)
Downloading pandas-2.2.2-cp311-cp311-macosx_11_0_arm64.whl (11.3 MB)
    11.3/11.3 MB
4.7 MB/s eta 0:00:0000:0100:01
Downloading numpy-2.0.1-cp311-cp311-macosx_14_0_arm64.whl (5.3 MB)
    5.3/5.3 MB
9.9 MB/s eta 0:00:00ta 0:00:01
Using cached pytz-2024.1-py2.py3-none-any.whl (505 kB)
Using cached tzdata-2024.1-py2.py3-none-any.whl (345 kB)
Installing collected packages: pytz, tzdata, numpy, pandas, fredapi
Successfully installed fredapi-0.5.2 numpy-2.0.1 pandas-2.2.2 pytz-2024.1
tzdata-2024.1
Note: you may need to restart the kernel to use updated packages.
```

```
[6]: import pandas as pd
import fredapi
```

```
[7]: api_key = '6b7f972da2d1e8428579d50c4ba98047'
fred = fredapi.Fred(api_key = api_key)
```

```
[8]: def gen_df(category, series):
    gen_ser = fred.get_series(series, frequency = 'q')
    return pd.DataFrame({'Date': gen_ser.index, category + ' : Billions of_
↳dollars': gen_ser.values})

def merge_dataframes(dataframes, on_column):
    merged_df = dataframes[0]
    for df in dataframes[1:]:
        merged_df = pd.merge(merged_df, df, on = on_column)
    return merged_df

dataframes_list = [
    gen_df('GDP', 'GDP'),
    gen_df('PCE', 'PCE'),
    gen_df('GPDI', 'GPDI'),
    gen_df('NETEXP', 'NETEXP'),
    gen_df('GovTotExp', 'W068RCQ027SBEA')
]

data = merge_dataframes(dataframes_list, 'Date')
data
```

```
[8]:      Date  GDP : Billions of dollars  PCE : Billions of dollars \
0   1960-01-01          542.648          326.4
1   1960-04-01          541.080          332.2
2   1960-07-01          545.604          332.1
3   1960-10-01          540.197          334.0
4   1961-01-01          545.018          334.5
..      ...
253 2023-04-01        27063.012        18419.0
254 2023-07-01        27610.128        18679.5
255 2023-10-01        27956.998        18914.5
256 2024-01-01        28269.174        19142.6
257 2024-04-01        28629.153        19377.7

      GPDI : Billions of dollars  NETEXP : Billions of dollars \
0          96.476          2.858
1          87.096          3.395
2          86.377          4.682
3          75.963          5.880
4          78.378          5.902
..      ...
253        4780.290        -806.093
254        4915.033        -779.231
```

255	4954.426	-783.734
256	5020.538	-834.896
257	5144.567	-894.362

	GovTotExp : Billions of dollars
0	144.233
1	147.417
2	150.459
3	153.780
4	157.254
..	...
253	9422.404
254	10007.677
255	9700.808
256	9925.034
257	10067.640

[258 rows x 6 columns]

```
[11]: %pip install matplotlib
```

```
Collecting matplotlib
  Downloading matplotlib-3.9.1-cp311-cp311-macosx_11_0_arm64.whl.metadata (11
kB)
Collecting contourpy>=1.0.1 (from matplotlib)
  Using cached contourpy-1.2.1-cp311-cp311-macosx_11_0_arm64.whl.metadata (5.8
kB)
Collecting cyclor>=0.10 (from matplotlib)
  Using cached cyclor-0.12.1-py3-none-any.whl.metadata (3.8 kB)
Collecting fonttools>=4.22.0 (from matplotlib)
  Downloading fonttools-4.53.1-cp311-cp311-macosx_11_0_arm64.whl.metadata (162
kB)
162.6/162.6
kB 5.9 MB/s eta 0:00:00
Collecting kiwisolver>=1.3.1 (from matplotlib)
  Using cached kiwisolver-1.4.5-cp311-cp311-macosx_11_0_arm64.whl.metadata (6.4
kB)
Requirement already satisfied: numpy>=1.23 in
/Users/neerajnamani/miniconda3/envs/py311/lib/python3.11/site-packages (from
matplotlib) (2.0.1)
Requirement already satisfied: packaging>=20.0 in
/Users/neerajnamani/miniconda3/envs/py311/lib/python3.11/site-packages (from
matplotlib) (23.2)
Collecting pillow>=8 (from matplotlib)
  Downloading pillow-10.4.0-cp311-cp311-macosx_11_0_arm64.whl.metadata (9.2 kB)
Collecting pyparsing>=2.3.1 (from matplotlib)
  Using cached pyparsing-3.1.2-py3-none-any.whl.metadata (5.1 kB)
```

Requirement already satisfied: python-dateutil>=2.7 in
 /Users/neerajnamani/miniconda3/envs/py311/lib/python3.11/site-packages (from
 matplotlib) (2.8.2)
 Requirement already satisfied: six>=1.5 in
 /Users/neerajnamani/miniconda3/envs/py311/lib/python3.11/site-packages (from
 python-dateutil>=2.7->matplotlib) (1.16.0)
 Downloading matplotlib-3.9.1-cp311-cp311-macosx_11_0_arm64.whl (7.8 MB)
 7.8/7.8 MB
 5.5 MB/s eta 0:00:0000:0100:01m
 Using cached contourpy-1.2.1-cp311-cp311-macosx_11_0_arm64.whl (245 kB)
 Using cached cycler-0.12.1-py3-none-any.whl (8.3 kB)
 Downloading fonttools-4.53.1-cp311-cp311-macosx_11_0_arm64.whl (2.2 MB)
 2.2/2.2 MB
 8.5 MB/s eta 0:00:00a 0:00:01
 Using cached kiwisolver-1.4.5-cp311-cp311-macosx_11_0_arm64.whl (66 kB)
 Downloading pillow-10.4.0-cp311-cp311-macosx_11_0_arm64.whl (3.4 MB)
 3.4/3.4 MB
 9.7 MB/s eta 0:00:00ta 0:00:01
 Using cached pyparsing-3.1.2-py3-none-any.whl (103 kB)
 Installing collected packages: pyparsing, pillow, kiwisolver, fonttools, cycler,
 contourpy, matplotlib
 Successfully installed contourpy-1.2.1 cycler-0.12.1 fonttools-4.53.1
 kiwisolver-1.4.5 matplotlib-3.9.1 pillow-10.4.0 pyparsing-3.1.2
 Note: you may need to restart the kernel to use updated packages.


```
[12]: # Visualization
import matplotlib.pyplot as plt
#separating date column from feature columns
date_column = 'Date'
feature_columns = data.columns.difference([date_column])
#set the plot
fig, ax = plt.subplots(figsize = (10,6))
fig.suptitle('Features vs Time', y = 1.02)

#graphing features onto plot
for i, feature in enumerate(feature_columns):
    ax.plot(data[date_column], data[feature], label = feature, color = plt.cm.
    viridis(i / len(feature_columns)))

#label axis
ax.set_xlabel('Date')
ax.set_ylabel('Billions of Dollars')
ax.legend(loc='upper left', bbox_to_anchor = (1,1))
#display plot
plt.show()
```

Matplotlib is building the font cache; this may take a moment.

Features vs Time



```
[16]: %pip install statsmodels
```

Collecting statsmodels

Downloading statsmodels-0.14.2-cp311-cp311-macosx_11_0_arm64.whl.metadata (9.2 kB)

Requirement already satisfied: numpy>=1.22.3 in /Users/neerajnamani/miniconda3/envs/py311/lib/python3.11/site-packages (from statsmodels) (2.0.1)

Requirement already satisfied: scipy!=1.9.2,>=1.8 in /Users/neerajnamani/miniconda3/envs/py311/lib/python3.11/site-packages (from statsmodels) (1.14.0)

Requirement already satisfied: pandas!=2.1.0,>=1.4 in /Users/neerajnamani/miniconda3/envs/py311/lib/python3.11/site-packages (from statsmodels) (2.2.2)

Collecting patsy>=0.5.6 (from statsmodels)

Downloading patsy-0.5.6-py2.py3-none-any.whl.metadata (3.5 kB)

Requirement already satisfied: packaging>=21.3 in /Users/neerajnamani/miniconda3/envs/py311/lib/python3.11/site-packages (from statsmodels) (23.2)

Requirement already satisfied: python-dateutil>=2.8.2 in /Users/neerajnamani/miniconda3/envs/py311/lib/python3.11/site-packages (from pandas!=2.1.0,>=1.4->statsmodels) (2.8.2)

Requirement already satisfied: pytz>=2020.1 in /Users/neerajnamani/miniconda3/envs/py311/lib/python3.11/site-packages (from pandas!=2.1.0,>=1.4->statsmodels) (2024.1)

Requirement already satisfied: tzdata>=2022.7 in /Users/neerajnamani/miniconda3/envs/py311/lib/python3.11/site-packages (from pandas!=2.1.0,>=1.4->statsmodels) (2024.1)

Requirement already satisfied: six in
 /Users/neerajnamani/miniconda3/envs/py311/lib/python3.11/site-packages (from
 patsy>=0.5.6->statsmodels) (1.16.0)
 Downloading statsmodels-0.14.2-cp311-macosx_11_0_arm64.whl (10.1 MB)
 10.1/10.1 MB
 15.4 MB/s eta 0:00:00 0:00:01
 Downloading patsy-0.5.6-py2.py3-none-any.whl (233 kB)
 233.9/233.9 kB
 17.2 MB/s eta 0:00:00
 Installing collected packages: patsy, statsmodels
 Successfully installed patsy-0.5.6 statsmodels-0.14.2
 Note: you may need to restart the kernel to use updated packages.

```
[17]: from statsmodels.tsa.stattools import adfuller
      #iterating through each feature
      for column in data.columns:
          if column!= 'Date':
              result = adfuller(data[column])
              print(f"ADF statistic for {column}: {result[0]}")
              print(f"P-value for {column}: {result[1]}")
              print("Critical Values:")
              for key, value in result[4].items():
                  print(f"{key}: {value}")
      # creating separation line between each feature
      print("\n" + "=" * 40 + "\n")
```

ADF statistic for GDP : Billions of dollars: 7.435656436752694
 P-value for GDP : Billions of dollars: 1.0
 Critical Values:
 1%: -3.4561550092339512
 5%: -2.8728972266578676
 10%: -2.5728222369384763

=====

ADF statistic for PCE : Billions of dollars: 6.8778373692584145
 P-value for PCE : Billions of dollars: 1.0
 Critical Values:
 1%: -3.4561550092339512
 5%: -2.8728972266578676
 10%: -2.5728222369384763

=====

ADF statistic for GPDI : Billions of dollars: 2.992625554816609
 P-value for GPDI : Billions of dollars: 1.0
 Critical Values:
 1%: -3.4560535712549925

5%: -2.8728527662442334
10%: -2.5727985212493754

=====

ADF statistic for NETEXP : Billions of dollars: 0.08039679949897652
P-value for NETEXP : Billions of dollars: 0.9646969244646914
Critical Values:
1%: -3.4564641849494113
5%: -2.873032730098417
10%: -2.572894516864816

=====

ADF statistic for GovTotExp : Billions of dollars: 2.7872131183908557
P-value for GovTotExp : Billions of dollars: 1.0
Critical Values:
1%: -3.4577787098622674
5%: -2.873608704758507
10%: -2.573201765981991

=====

```
[26]: # differencing and storing original dataset
data_diff = data.drop('Date', axis = 1).diff().dropna()
#printing ADF test for new dataset
for column in data_diff.columns:
    result = adfuller(data_diff[column])
    print(f"ADF Statistic for {column}: {result[0]}")
    print(f"P-value for {column}: {result[1]}")
    print("Critical Values:")
    for key, value in result[4].items():
        print(f" {key}: {value}")

    print("\n" + "=" * 40 + "\n")
```

ADF Statistic for GDP : Billions of dollars: -3.886345364333152
P-value for GDP : Billions of dollars: 0.002137667738212588
Critical Values:
1%: -3.4565688966099373
5%: -2.8730786194395455
10%: -2.5729189953388762

=====

ADF Statistic for PCE : Billions of dollars: -6.039053627599207
P-value for PCE : Billions of dollars: 1.3598172420763878e-07

Critical Values:

1%: -3.456360306409983
5%: -2.8729872043802356
10%: -2.572870232500465

=====

ADF Statistic for GPD I : Billions of dollars: -14.608670267780347

P-value for GPD I : Billions of dollars: 4.0558964179422356e-27

Critical Values:

1%: -3.4561550092339512
5%: -2.8728972266578676
10%: -2.5728222369384763

=====

ADF Statistic for NETEXP : Billions of dollars: -8.804306346787321

P-value for NETEXP : Billions of dollars: 2.071671900401648e-14

Critical Values:

1%: -3.4564641849494113
5%: -2.873032730098417
10%: -2.572894516864816

=====

ADF Statistic for GovTotExp : Billions of dollars: -3.108588974313343

P-value for GovTotExp : Billions of dollars: 0.025924060990244243

Critical Values:

1%: -3.457664132155201
5%: -2.8735585105960224
10%: -2.5731749894132916

=====

```
[30]: from statsmodels.tsa.stattools import grangercausalitytests
lags = [6,9,1,1]

for column, lag in zip(data_diff.columns, lags):
    df_new = data_diff[['GDP : Billions of dollars', column]]
    print(f'For: {column}')
    gc_res = grangercausalitytests(df_new, lag)
    print("\n" + "=" * 40 + "\n")
```

For: GDP : Billions of dollars

Granger Causality

number of lags (no zero) 1

ssr based F test: F=-0.0000 , p=1.0000 , df_denom=254, df_num=1
 ssr based chi2 test: chi2=-0.0000 , p=1.0000 , df=1
 likelihood ratio test: chi2=-0.0000 , p=1.0000 , df=1
 parameter F test: F=2.1900 , p=0.1401 , df_denom=254, df_num=1

Granger Causality

number of lags (no zero) 2

ssr based F test: F=0.0000 , p=1.0000 , df_denom=252, df_num=2
 ssr based chi2 test: chi2=0.0000 , p=1.0000 , df=2
 likelihood ratio test: chi2=-0.0000 , p=1.0000 , df=2
 parameter F test: F=7.8050 , p=0.0005 , df_denom=252, df_num=2

Granger Causality

number of lags (no zero) 3

ssr based F test: F=-0.0000 , p=1.0000 , df_denom=250, df_num=3
 ssr based chi2 test: chi2=-0.0000 , p=1.0000 , df=3
 likelihood ratio test: chi2=-0.0000 , p=1.0000 , df=3
 parameter F test: F=8.7149 , p=0.0000 , df_denom=250, df_num=3

Granger Causality

number of lags (no zero) 4

ssr based F test: F=0.0000 , p=1.0000 , df_denom=248, df_num=4
 ssr based chi2 test: chi2=0.0000 , p=1.0000 , df=4
 likelihood ratio test: chi2=-0.0000 , p=1.0000 , df=4
 parameter F test: F=7.2269 , p=0.0000 , df_denom=248, df_num=4

Granger Causality

number of lags (no zero) 5

ssr based F test: F=0.0000 , p=1.0000 , df_denom=246, df_num=5
 ssr based chi2 test: chi2=0.0000 , p=1.0000 , df=5
 likelihood ratio test: chi2=-0.0000 , p=1.0000 , df=5
 parameter F test: F=7.2337 , p=0.0000 , df_denom=246, df_num=5

Granger Causality

number of lags (no zero) 6

ssr based F test: F=0.0000 , p=1.0000 , df_denom=244, df_num=6
 ssr based chi2 test: chi2=0.0000 , p=1.0000 , df=6
 likelihood ratio test: chi2=-0.0000 , p=1.0000 , df=6
 parameter F test: F=5.9489 , p=0.0000 , df_denom=244, df_num=6

=====

For: PCE : Billions of dollars

Granger Causality

number of lags (no zero) 1

ssr based F test: F=0.9211 , p=0.3381 , df_denom=253, df_num=1
 ssr based chi2 test: chi2=0.9320 , p=0.3343 , df=1

likelihood ratio test: $\chi^2=0.9303$, $p=0.3348$, $df=1$
parameter F test: $F=0.9211$, $p=0.3381$, $df_{denom}=253$, $df_{num}=1$

Granger Causality

number of lags (no zero) 2

ssr based F test: $F=0.7615$, $p=0.4680$, $df_{denom}=250$, $df_{num}=2$
ssr based χ^2 test: $\chi^2=1.5534$, $p=0.4599$, $df=2$
likelihood ratio test: $\chi^2=1.5487$, $p=0.4610$, $df=2$
parameter F test: $F=0.7615$, $p=0.4680$, $df_{denom}=250$, $df_{num}=2$

Granger Causality

number of lags (no zero) 3

ssr based F test: $F=0.8279$, $p=0.4796$, $df_{denom}=247$, $df_{num}=3$
ssr based χ^2 test: $\chi^2=2.5542$, $p=0.4656$, $df=3$
likelihood ratio test: $\chi^2=2.5414$, $p=0.4679$, $df=3$
parameter F test: $F=0.8279$, $p=0.4796$, $df_{denom}=247$, $df_{num}=3$

Granger Causality

number of lags (no zero) 4

ssr based F test: $F=0.8084$, $p=0.5208$, $df_{denom}=244$, $df_{num}=4$
ssr based χ^2 test: $\chi^2=3.3530$, $p=0.5006$, $df=4$
likelihood ratio test: $\chi^2=3.3310$, $p=0.5040$, $df=4$
parameter F test: $F=0.8084$, $p=0.5208$, $df_{denom}=244$, $df_{num}=4$

Granger Causality

number of lags (no zero) 5

ssr based F test: $F=1.7989$, $p=0.1137$, $df_{denom}=241$, $df_{num}=5$
ssr based χ^2 test: $\chi^2=9.4049$, $p=0.0940$, $df=5$
likelihood ratio test: $\chi^2=9.2337$, $p=0.1001$, $df=5$
parameter F test: $F=1.7989$, $p=0.1137$, $df_{denom}=241$, $df_{num}=5$

Granger Causality

number of lags (no zero) 6

ssr based F test: $F=2.0568$, $p=0.0591$, $df_{denom}=238$, $df_{num}=6$
ssr based χ^2 test: $\chi^2=13.0150$, $p=0.0428$, $df=6$
likelihood ratio test: $\chi^2=12.6888$, $p=0.0483$, $df=6$
parameter F test: $F=2.0568$, $p=0.0591$, $df_{denom}=238$, $df_{num}=6$

Granger Causality

number of lags (no zero) 7

ssr based F test: $F=1.8111$, $p=0.0859$, $df_{denom}=235$, $df_{num}=7$
ssr based χ^2 test: $\chi^2=13.4870$, $p=0.0611$, $df=7$
likelihood ratio test: $\chi^2=13.1358$, $p=0.0689$, $df=7$
parameter F test: $F=1.8111$, $p=0.0859$, $df_{denom}=235$, $df_{num}=7$

Granger Causality

number of lags (no zero) 8

ssr based F test: $F=1.6846$, $p=0.1030$, $df_{denom}=232$, $df_{num}=8$

```

ssr based chi2 test:   chi2=14.4640 , p=0.0704 , df=8
likelihood ratio test: chi2=14.0595 , p=0.0802 , df=8
parameter F test:      F=1.6846 , p=0.1030 , df_denom=232, df_num=8

```

Granger Causality

number of lags (no zero) 9

```

ssr based F test:      F=1.5375 , p=0.1358 , df_denom=229, df_num=9
ssr based chi2 test:   chi2=14.9857 , p=0.0913 , df=9
likelihood ratio test: chi2=14.5504 , p=0.1040 , df=9
parameter F test:      F=1.5375 , p=0.1358 , df_denom=229, df_num=9

```

=====

For: GPDI : Billions of dollars

Granger Causality

number of lags (no zero) 1

```

ssr based F test:      F=0.4007 , p=0.5273 , df_denom=253, df_num=1
ssr based chi2 test:   chi2=0.4055 , p=0.5243 , df=1
likelihood ratio test: chi2=0.4052 , p=0.5244 , df=1
parameter F test:      F=0.4007 , p=0.5273 , df_denom=253, df_num=1

```

=====

For: NETEXP : Billions of dollars

Granger Causality

number of lags (no zero) 1

```

ssr based F test:      F=9.6776 , p=0.0021 , df_denom=253, df_num=1
ssr based chi2 test:   chi2=9.7923 , p=0.0018 , df=1
likelihood ratio test: chi2=9.6097 , p=0.0019 , df=1
parameter F test:      F=9.6776 , p=0.0021 , df_denom=253, df_num=1

```

=====

```

[33]: # Ensure the column name is correct and consistent
split_index = int(len(data_diff) * 0.90)
train_data = data_diff.iloc[:split_index]
test_data = data_diff.iloc[split_index:]

# Assigning GDP column to target variable
X_train = train_data.drop('GDP : Billions of dollars', axis=1)
y_train = train_data['GDP : Billions of dollars']
X_test = test_data.drop('GDP : Billions of dollars', axis=1)
y_test = test_data['GDP : Billions of dollars']

```

```
[36]: %pip install scikit-learn
```

```
Collecting scikit-learn
  Downloading scikit_learn-1.5.1-cp311-cp311-macosx_12_0_arm64.whl.metadata (12
kB)
Requirement already satisfied: numpy>=1.19.5 in
/Users/neerajnamani/miniconda3/envs/py311/lib/python3.11/site-packages (from
scikit-learn) (2.0.1)
Requirement already satisfied: scipy>=1.6.0 in
/Users/neerajnamani/miniconda3/envs/py311/lib/python3.11/site-packages (from
scikit-learn) (1.14.0)
Collecting joblib>=1.2.0 (from scikit-learn)
  Using cached joblib-1.4.2-py3-none-any.whl.metadata (5.4 kB)
Collecting threadpoolctl>=3.1.0 (from scikit-learn)
  Using cached threadpoolctl-3.5.0-py3-none-any.whl.metadata (13 kB)
Downloading scikit_learn-1.5.1-cp311-cp311-macosx_12_0_arm64.whl (11.0 MB)
    11.0/11.0 MB
7.6 MB/s eta 0:00:0000:0100:01
Using cached joblib-1.4.2-py3-none-any.whl (301 kB)
Using cached threadpoolctl-3.5.0-py3-none-any.whl (18 kB)
Installing collected packages: threadpoolctl, joblib, scikit-learn
Successfully installed joblib-1.4.2 scikit-learn-1.5.1 threadpoolctl-3.5.0
Note: you may need to restart the kernel to use updated packages.
```

```
[38]: from sklearn.ensemble import RandomForestRegressor

rf_model = RandomForestRegressor(n_estimators = 100, random_state = 42)
rf_model.fit(X_train, y_train)

y_pred = rf_model.predict(X_test)

import matplotlib.pyplot as plt
from sklearn.metrics import mean_squared_error, r2_score

# Define the printevals function
def printevals(y_true, y_pred):
    mse = mean_squared_error(y_true, y_pred)
    rmse = mse ** 0.5
    r2 = r2_score(y_true, y_pred)
    print(f"Mean Squared Error: {mse}")
    print(f"Root Mean Squared Error: {rmse}")
    print(f"R^2 Score: {r2}")

# Define the plotresults function
def plotresults(title):
    plt.figure(figsize=(10, 5))
    plt.plot(y_test.reset_index(drop=True), label='Actual GDP')
    plt.plot(y_pred, label='Forecasted GDP', linestyle='--')
```

```

plt.title(title)
plt.xlabel('Time')
plt.ylabel('GDP : Billions of dollars')
plt.legend()
plt.show()

# Assuming you have already trained your model and made predictions
# rf_model = ... # Your Random Forest model
# y_pred = rf_model.predict(X_test)


# Evaluate and plot results
printevals(y_test, y_pred)
plotresults('Actual vs Forecasted GDP using Random Forest')

```

Mean Squared Error: 232807.4460297813

Root Mean Squared Error: 482.50123940750797

R² Score: 0.18876026061487017



```

[39]: from sklearn.neighbors import KNeighborsRegressor
# iterate over all k=1 to k=10
for i in range(1,10):
    knn_model = KNeighborsRegressor(n_neighbors=i)
    knn_model.fit(X_train, y_train)

    y_pred = knn_model.predict(X_test)
    print(f'for k = {i}')
    printevals(y_test,y_pred)

```

```
print("\n" + "=" * 40 + "\n")
```

```
for k = 1
Mean Squared Error: 240281.69660884605
Root Mean Squared Error: 490.1853696397375
R^2 Score: 0.16271552194665762

=====

for k = 2
Mean Squared Error: 254614.91029419223
Root Mean Squared Error: 504.59380722933196
R^2 Score: 0.1127700724649251

=====

for k = 3
Mean Squared Error: 258879.7438215
Root Mean Squared Error: 508.80226397049375
R^2 Score: 0.0979088534695004

=====

for k = 4
Mean Squared Error: 257326.62522003113
Root Mean Squared Error: 507.27371824295324
R^2 Score: 0.10332084329618574

=====

for k = 5
Mean Squared Error: 266163.3469519337
Root Mean Squared Error: 515.9102121027782
R^2 Score: 0.07252844401059544

=====

for k = 6
Mean Squared Error: 271011.1762016729
Root Mean Squared Error: 520.5873377269878
R^2 Score: 0.05563572084297408

=====

for k = 7
Mean Squared Error: 270831.94846411614
Root Mean Squared Error: 520.4151693255262
```

R² Score: 0.056260256980394496

=====

for k = 8

Mean Squared Error: 275173.34869414533

Root Mean Squared Error: 524.5696795413793

R² Score: 0.04113223401018018

=====

for k = 9

Mean Squared Error: 275421.8648684799

Root Mean Squared Error: 524.8065023115471

R² Score: 0.04026625570949305

=====

```
[40]: #applying model with optimal k values
knn_model = KNeighborsRegressor(n_neighbors = 2)
knn_model.fit(X_train, y_train)

y_pred = knn_model.predict(X_test)

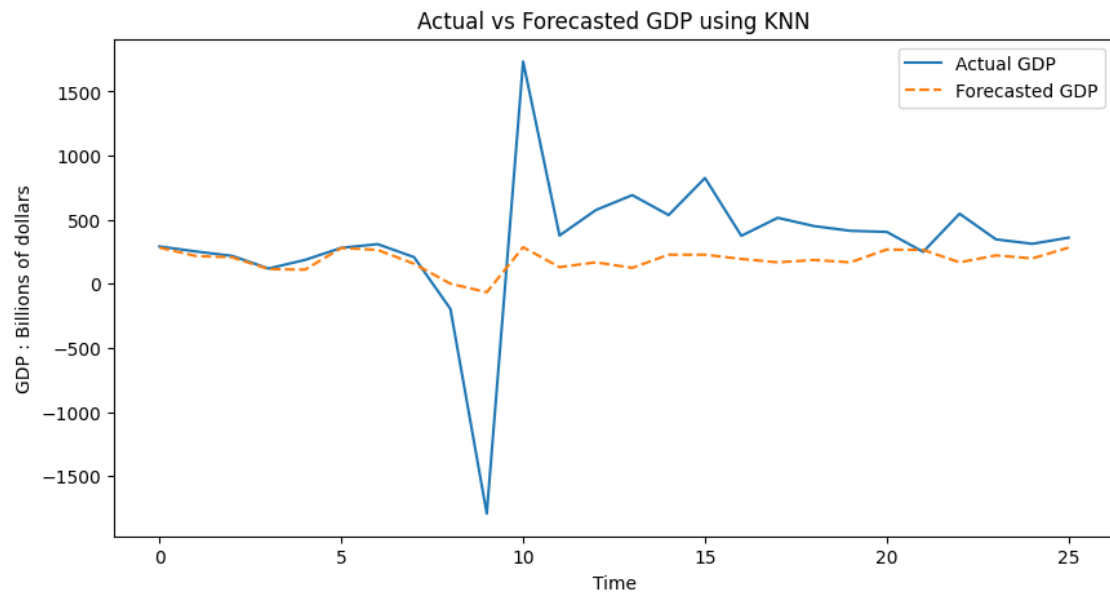
printevals(y_test, y_pred)

plotresults('Actual vs Forecasted GDP using KNN')
```

Mean Squared Error: 254614.91029419223

Root Mean Squared Error: 504.59380722933196

R² Score: 0.1127700724649251



[]: